

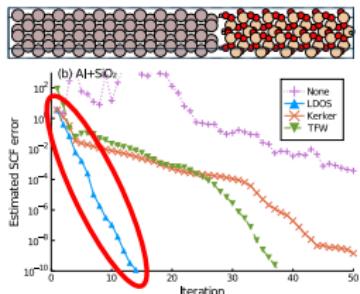
DFTK.jl: An introduction to a multidisciplinary electronic-structure code

Michael F. Herbst

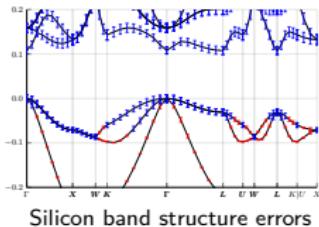
Mathematics for Materials Modelling (matmat.org), EPFL

21 October 2024

https://michael-herbst.com/talks/2024.10.21_JuliaMolSim_DFTK.pdf



Novel materials simulation algorithms



Silicon band structure errors

Energy consumption of materials discovery



- Current solutions limited by properties of available materials
 ⇒ Innovation driven by **discovering new materials**
 - **Experimental** research extremely **energy intensive**
 - 1 fume hood \simeq 2-3 average households¹
- ⇒ Complement experiment by **computational materials discovery**

¹D. Wesolowski *et. al.* Int. J. Sustain. High. Edu. 11, 217 (2010).

High-throughput materials screening



$$\min_{\Psi} \langle \Psi, H\Psi \rangle$$



- Energy consumption ?



High-throughput materials screening



$$\min_{\Psi} \langle \Psi, H\Psi \rangle$$

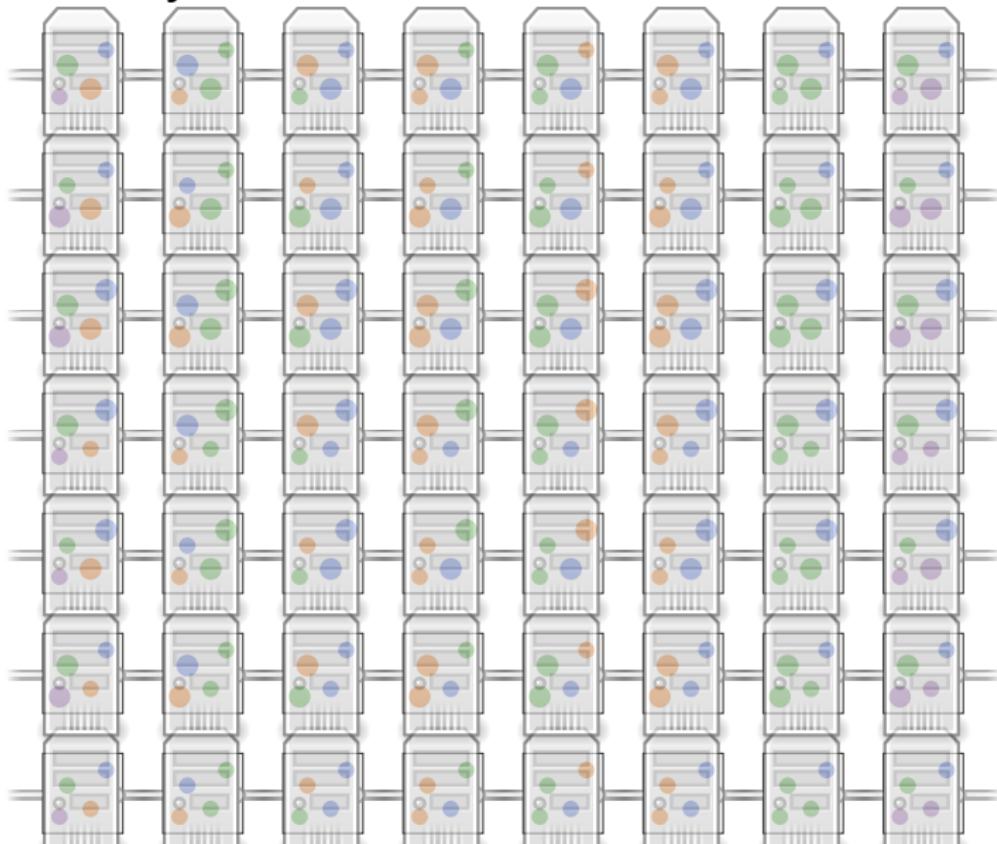


- Energy consumption ?
 - 8h of 36-core processor
 \simeq 4h of average household
 \simeq 1 CHF

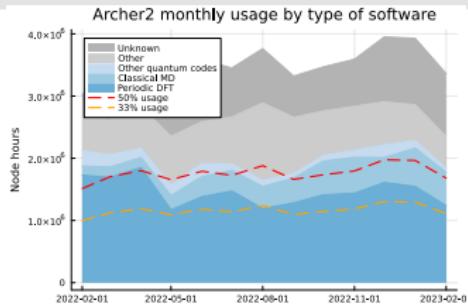
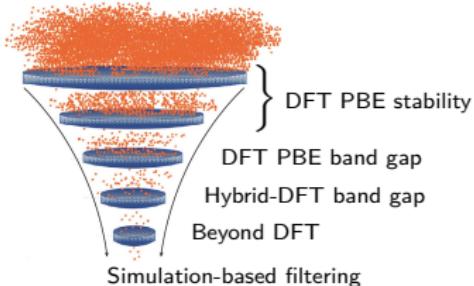


High-throughput materials screening

- We can **fully automate** this !

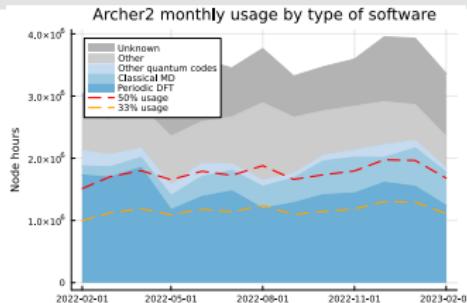
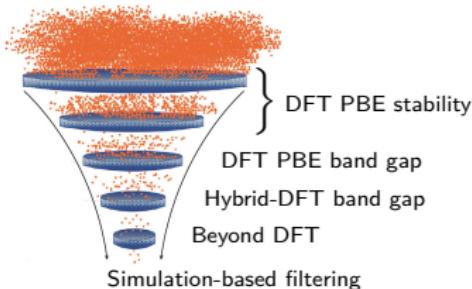


Computational materials discovery



- Goal: Only promising candidates made in the lab
- Systematic simulations on $\simeq 10^4 - 10^6$ compounds
 - Noteworthy share of world's supercomputing resources

Computational materials discovery



- **Goal:** Only promising candidates made in the lab
- Systematic simulations on $\simeq 10^4 - 10^6$ compounds
 - Noteworthy share of world's supercomputing resources
- Energy consumption of LUMI (one of the most efficient):
 - 60 million kWh / year $\simeq 1.5$ EPFLs $\simeq 14000$ households



Challenges of high-throughput regime



AFLOW
Automatic FLOW for Materials Discovery



MATERIALS CLOUD

AiiDA



- Complexity of multiscale materials modelling
 - Many parameters to choose (algorithms, tolerances, models)
 - Automated workflows & data management software (see above)
- Despite elaborate heuristics: Thousands of failed calculations
 - ⇒ Wasted resources
 - ⇒ Increased human attention (limits throughput)
- Traversing the design space
 - How to best optimise material properties
 - How much accuracy is needed ?
 - How could we explore unusual gradients ?

A focus on robust materials simulations

- Goal in ~~Mat~~ group:

- Obtain **reliable & efficient** simulations
- Develop and employ mathematically sound **error indicators**
- Transform **empirical wisdom** to built-in **convergence guarantees**

⇒ Understand **where and how** to spend efforts best

- Practical error indicators:

- Automatic & robust verification
- Multi-fidelity statistical surrogates
- Active learning of missing physics

- Leverage inexactness:

- Error balancing: Optimal adaptive parameter selection
- Adaptive tolerances & selective precision

⇒ Multidisciplinary expertise required

A focus on robust materials simulations

- Goal in ~~Mat~~ group:

- Obtain reliable & efficient simulations
- Develop and employ mathematically sound error indicators
- Transform empirical wisdom to built-in convergence guarantees

⇒ Understand where and how to spend efforts best

- Practical error indicators:

- Automatic & robust verification
- Multi-fidelity statistical surrogates
- Active learning of missing physics

- Leverage inexactness:

- Error balancing: Optimal adaptive parameter selection
- Adaptive tolerances & selective precision

⇒ Multidisciplinary expertise required

Difficulties of cross-disciplinary research

(A computational science point of view . . .)

- Community conventions . . .
 - Language barriers, publication culture, speed of research, . . .
- . . . that are cemented in software:
 - Priorities differ ⇒ What is considered “a good code” differs

Mathematical software

- **Goal:** Numerical experiments
- **Scope:** Reduced models
- **High-level language:**
Matlab, python, . . .
- **Lifetime:** 1 paper
- **Size:** < 1k lines
- Does not care about performance

Application software

- **Goal:** Modelling physics
- **Scope:** All relevant systems
- **Mix of languages:**
C, FORTRAN, python, . . .
- **Lifetime:** 100 manyears
- **Size:** 100k – 1M lines
- Obliged to write performant code

Difficulties of cross-disciplinary research

(A computational science point of view . . .)

Mathematical software

- **Goal:** Numerical experiments
- **Scope:** Reduced models
- **High-level language:**
Matlab, python, . . .
- **Lifetime:** 1 paper
- **Size:** < 1k lines
- Does not care about performance

Application software

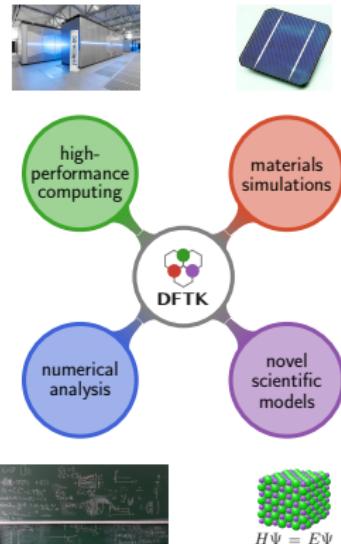
- **Goal:** Modelling physics
- **Scope:** All relevant systems
- **Mix of languages:**
C, FORTRAN, python, . . .
- **Lifetime:** 100 manyears
- **Size:** 100k – 1M lines
- Obliged to write performant code

- Working with these codes requires different skillsets
⇒ Orthogonal developer & user communities

- Obstacle for knowledge transfer:
 - Mathematical methods never tried in practical setting
(and may well not work well in the real world)
 - Some issues cannot be studied with mathematical codes
(and mathematicians may never get to know of them)

- **Hypothesis:** People compose if software composes

Density-functional toolkit¹ — <https://dftk.org>



- Julia code for cross-disciplinary research:
 - Allows restriction to relevant model problems,
 - and scale-up to application regime (1000 electrons)
 - Sizeable feature set in 7500 lines of code
 - Norm-conserving pseudos, mGGA functionals, response
 - Integrated with high-throughput: MARVEL AiiDA
- Fully composable due to Julia abstractions:
 - Arbitrary precision (32bit, >64bit, ...)
 - Algorithmic differentiation (AD)
 - HPC tools: GPU acceleration, MPI parallelisation
- Accessible high-productivity research framework:
 - Key contributions by undergrads (AD, GPU, Pseudos, ...)
 - Over 30 contributors in 5 years (Maths, physics, CS, ...)

¹MFH, A. Levitt, E. Cancès. JuliaCon Proc. 3, 69 (2021).

Contents

1 Example use cases for  DFTK

2 API and ecosystem integration

Density-functional theory (insulators)

- Goal: Understand electronic structures (Many-body quantum system)
- DFT approximation: Effective single-particle model

$$\left\{ \begin{array}{l} \forall i \in 1 \dots N : \left(-\frac{1}{2} \Delta + V(\rho_\Phi) \right) \psi_i = \varepsilon_i \psi_i, \\ V(\rho) = V_{\text{ext}} + V_{\text{Hxc}}(\rho), \\ \rho_\Phi = \sum_{i=1}^N |\psi_i|^2, \end{array} \right.$$

- Self-consistent field (SCF) fixed-point problem

$$\rho(V(\rho)) = \rho$$

- Density mixing (preconditioner P , damping α)

$$\rho_{n+1} = \rho_n + \alpha P^{-1} [\rho(V(\rho_n)) - \rho_n]$$

- Best P & α highly system dependent (metal, insulator, ...)

- Usually chosen by trial and error (Impact on energy consumption ...)

Self-consistent field problem

- Density-mixing SCF procedure (preconditioner P , damping α)

$$\rho_{n+1} = \rho_n + \alpha P^{-1} [\rho(V(\rho_n)) - \rho_n]$$

- Near a fixed-point the error goes as

$$e_{n+1} \simeq [1 - \alpha P^{-1} \varepsilon^\dagger] e_n$$

with dielectric matrix $\varepsilon^\dagger = (1 - \chi_0 K)$, $K(\rho) = V'(\rho)$, $\chi_0(V) = \rho'(V)$

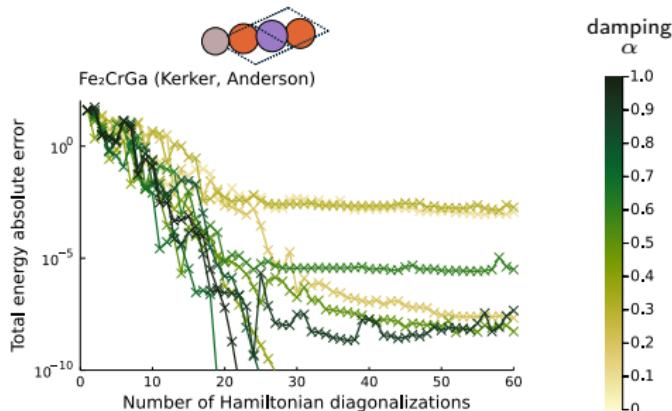
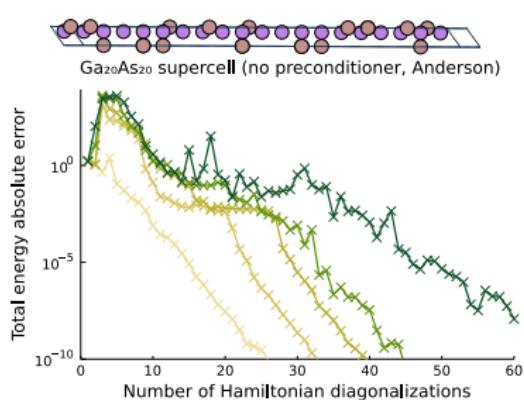
- Convergence iff $-1 < [1 - \alpha P^{-1} \varepsilon^\dagger] < 1$

- Dielectric matrix ε^\dagger : Depends on physics (conduction, screening)
- By second-order conditions: $\varepsilon^\dagger \geq 0$ (near fixed point)

⇒ Ideal preconditioner has $P^{-1} \varepsilon^\dagger \approx I$

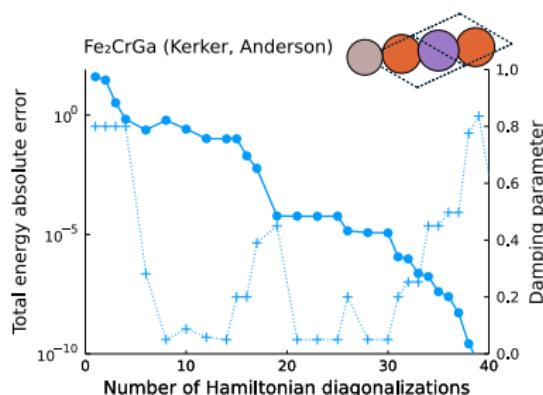
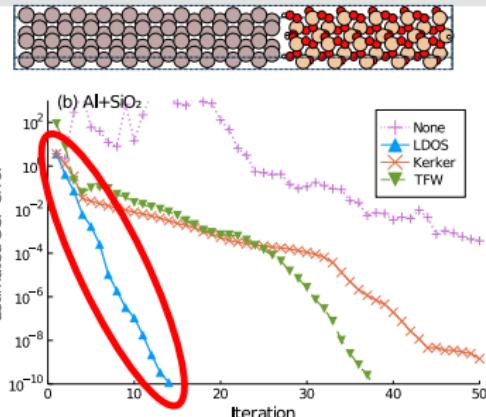
- Note: P needs to adapt to physics of unknown system!
- No such P available: Choose α appropriately (Trial and error)

Illustration: Guessing a suitable damping α can be hard



- Inefficient standard damping ($0.6 - 0.8$)
- Surprisingly small damping for smooth convergence
- Heusler alloy: Materials class with unusual magnetic properties
⇒ Numerically challenging behaviour
- SCF irregular: α versus convergence
- Usual heuristics breaks:
Larger damping is better

Self-adapting black-box algorithms



- Preconditioning inhomogeneous systems (surfaces, clusters, ...)
 - LDOS preconditioner¹: Parameter-free and self-adapting
 - ca. 50% less iterations
- Damping α adapted in each step (using tailored quadratic model)
 - Avoids trial and error (but may have a small overhead)
 - Safeguard with theoretical guarantees²

⇒ Maths / physics collaboration:

Exchange of ideas between simplified & practical settings crucial

¹MFH, A. Levitt. J. Phys. Condens. Matter **33**, 085503 (2021).

²MFH, A. Levitt. J. Comput. Phys. **459**, 111127 (2022).

Response, properties and algorithmic differentiation

- DFT properties: Response of system to external changes:

- Connection Theory \Leftrightarrow Experiment

- Modelling: Potential $V(\theta, \rho)$ depends on parameters θ
(e.g. atomic positions, el. field)

- SCF procedure yields fixed-point density ρ_*

$$0 = \rho(V(\theta, \rho_*)) - \rho_*$$

\Rightarrow Defines implicit function $\rho_*(\theta)$

- Properties are derivatives:

- Forces (energy wrt. position), dipole moment (energy wrt. el. field), elasticity (energy cross-response to lattice deformation), phonons, electronic spectra, ...

\Rightarrow Great application for algorithmic differentiation !

- Byproduct: Arbitrary derivatives

- Sensitivities, improved training of surrogates ...

AD for stresses keeps code accessible

Stress =

$$\frac{1}{\det(\mathbf{L})} \left. \frac{\partial E[P_*, (\mathbf{I} + \mathbf{M}) \mathbf{L}]}{\partial \mathbf{M}} \right|_{\mathbf{M}=0}$$

```
# Run SCF, get P*
scfres = self_consistent_field(basis)
L = basis.model.lattice
stress = 1/det(L) * gradient(
    M -> recompute_energy(
        scfres, (I + M) * L),
    zero(L)
)
```

- Stress computation (Definition vs.  code)¹
- Post-processing step \Rightarrow Not performance critical
- Comparison of implementation complexity:
 -  **DFTK**: 20 lines¹ (forward-mode algorithmic differentiation)
 - Quantum-Espresso: 1700 lines²
 - Initial version: \simeq 10-week GSoC project

¹<https://github.com/JuliaMolSim/DFTK.jl/blob/master/src/postprocess/stresses.jl>

²<https://github.com/QEF/q-e/blob/develop/PW/src>

Arbitrary derivatives: Need efficient response

- Full DFT equivalent is density-functional perturbation theory

$$\frac{\partial \rho_*}{\partial \theta} = [1 - \chi_0 K]^{-1} \chi_0 \frac{\partial V}{\partial \theta} \quad (2)$$

- Challenge: Need *many* applications of χ_0 :
 - Each requires solving N Sternheimer equations

$$(\tilde{H} - \varepsilon_i) \delta\psi_i = -P \delta V \psi_i \quad \forall i = 1, \dots, N$$

$H = -\frac{1}{2}\Delta + V$, $\tilde{H} = P H P$ and P some projector (ε_i, ψ_i) eigenpairs of H

⇒ Nested iterative problem . . . which can be ill-conditioned

Sternheimer equations

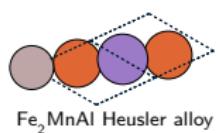
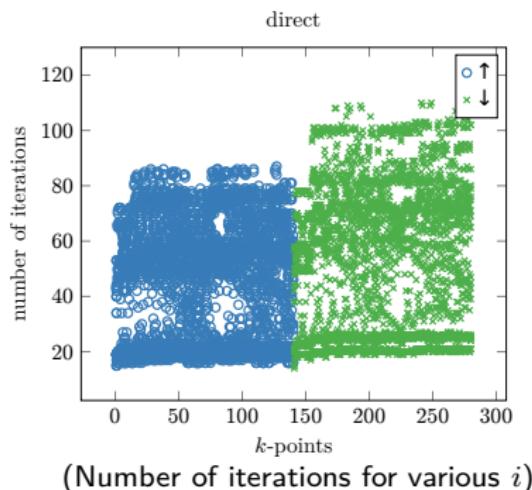
- Product $\chi_0 \delta V$ requires solving Sternheimer equations

$$(\tilde{H} - \varepsilon_i) \delta\psi_i = -P \delta V \psi_i \quad \forall i = 1, \dots, N$$

$H = -\frac{1}{2}\Delta + V$, $\tilde{H} = PHP$ and P some projector

(ε_i, ψ_i) eigenpairs of H

- ⇒ Badly conditioned for metallic systems (ε_i near eigenvalue of \tilde{H})



Schur complement approach to response¹

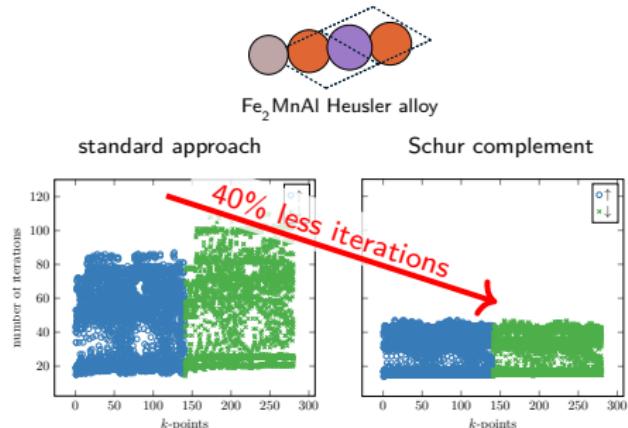
- Numerics of eigensolver:
We have N_{ex} “extra” bands
- Use these to partition \tilde{H} :

$$\tilde{H} = \begin{pmatrix} E_{\text{ex}} & \mathbf{C} \\ \mathbf{C}^\dagger & \mathbf{R} \end{pmatrix}$$

$E_{\text{ex}} = \text{diag}(\varepsilon_{N+1}, \dots, \varepsilon_{N+N_{\text{ex}}})$
& \mathbf{C} , \mathbf{R} projections of \tilde{H}

⇒ Use **Schur complement**:
Better-conditioned systems

$$(\mathbf{R} - \varepsilon_i)x = b$$

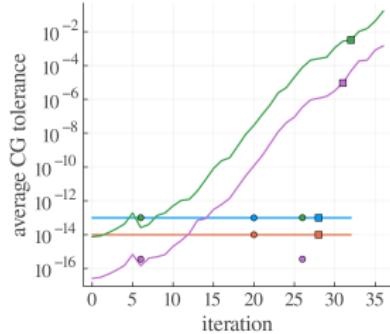
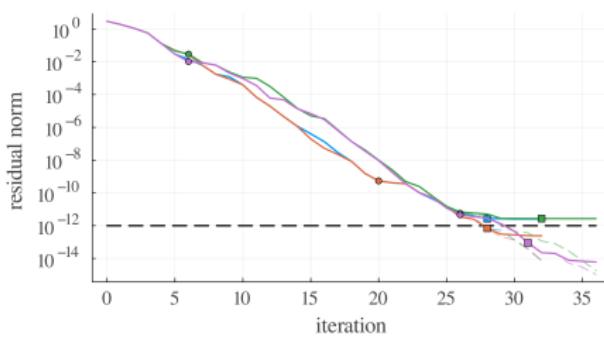


- Schur-based approach tames CG
- ca. 40% less iterations
- Development guided using a “real material”

¹E. Cancès, MFH, G. Kemlin, et. al. Lett. Math. Phys. **113**, 21 (2023).

WIP: Inexact Krylov methods

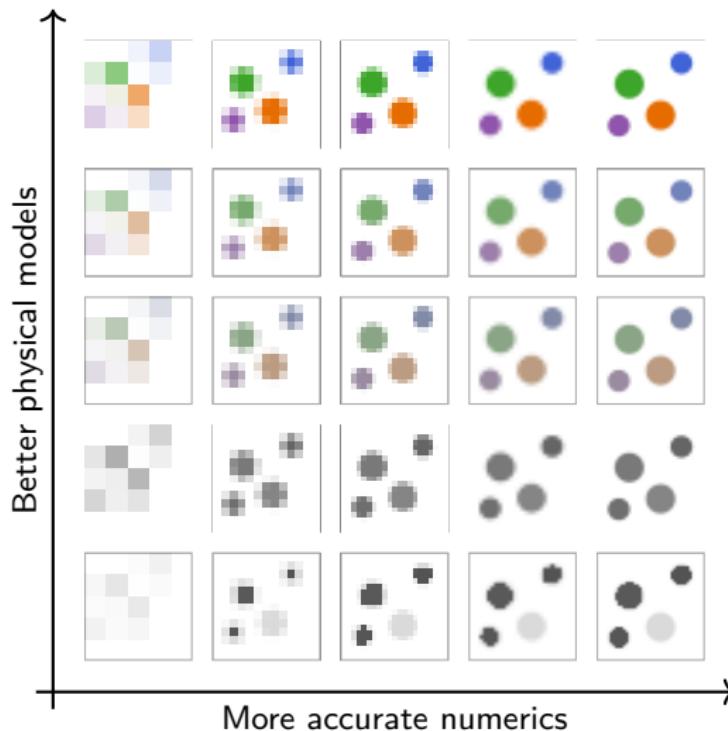
- DFPT + Sternheimer: Nested linear problems
- Inexact Krylov methods:¹ Framework to tolerate *less tight* solutions of Sternheimer
- First results indicate 25%–50% less Hamiltonian applications (the expensive step)



Bonan Sun

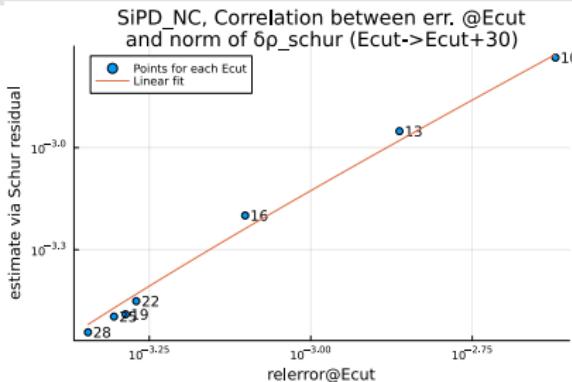
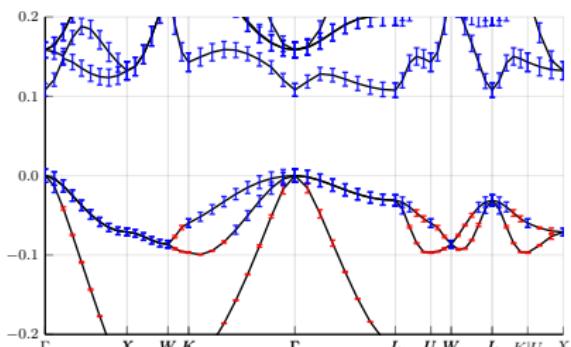
¹V. Simoncini, D. Szyld. SIAM J. Sci. Comput., **25**, 454 (2003).

Case for error control: Error comes in different flavours



- Ideally want to **balance** errors
- ⇒ Need reliable error indicators !

Numerical error: Analytical techniques

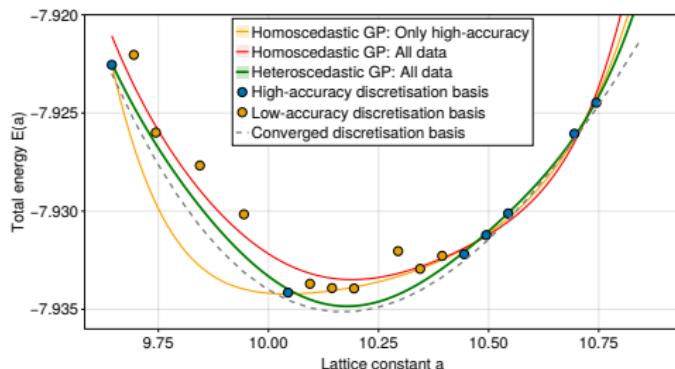


¹MFH, A. Levitt, E. Cancès. Faraday Discuss. **223**, 227 (2020).

²E. Cancès, G. Dusson, G. Kemlin et al. SIAM J. Sci. Comp., **44**, B1312 (2022).

³E. Cancès, G. Kemlin, A. Levitt. J. Matrix Anal. Appl., **42**, 243 (2021).

WIP: Heteroscedastic regression models



- 1D proof of principle: energy-volume curve (Equation of state)
- High-dimensional regression problems: Data is scarce
- Error δ_i can be estimated \Rightarrow supply to GP
- For example: Heteroscedastic model:

$$E_i = \text{DFT}(a_i) + \varepsilon_i \quad \varepsilon_i \sim \mathcal{N}(0, \delta_i)$$



Anna Paulish

DFT error: Computing model sensitivities

- Consider model sensitivity of force $\mathcal{F}(\rho_*(\theta))$:

$$\frac{d\mathcal{F}}{d\theta} = \frac{\partial \mathcal{F}}{\partial \rho_{\text{SCF}}} \frac{\partial \rho_*}{\partial \theta} \quad (1)$$

- Computed by response theory (we've seen this before !):

$$\frac{\partial \rho_*}{\partial \theta} = [1 - \chi_0 K]^{-1} \chi_0 \frac{\partial V}{\partial \theta}$$

- Parameters appear in innermost layer (model definition)
 - Each DFT model: Different derivatives $\frac{\partial V}{\partial \theta}$ (can be horrible)
 - Each quantity of interest: Different sensitivity expression (1)
⇒ Combinatorial explosion

WIP: Sensitivity analysis in one line of code

-  **DFTK**: Algorithmic differentiation (AD)
 - Generic framework for derivatives: Request gradient, AD delivers
 - ⇒ New properties/derivatives by non-DFT experts!
- ⇒ Setting for uncertainty quantification:
 - Pseudopotential sensitivity of electronic density



Niklas Schmitz

High-level structure of density-functional theory

- Energy minimisation problem (discretised setting):

$$\min_{D \in \mathcal{P}} E(D) = \min_{D \in \mathcal{P}} [\text{tr}(H_0 D) + E_{\text{Hxc}}(D)]$$

- Non-linear, non-convex Riemannian optimisation (\mathcal{P} : Grassmannian)

- What we care about: Illustration on model problem $x_* = \min_{x \in \mathbb{R}^N} E(x)$

- Numerical methods: $x_{k+1} = x_k - \alpha \nabla E(x_k)$ (SCF, direct min.)

- Convergence depends on $1 - \alpha \nabla^2 E(x_*)$
- Need to understand $\nabla^2 E(x_*)$ for preconditioning

- Response, properties and algorithmic differentiation

- Solution to $\min_x E(x, \theta)$ satisfies

$$x(\theta) \approx x_* - \theta \nabla^2 E(x_*, 0)^{-1} \frac{\partial}{\partial \theta} \nabla E(x_*, 0)$$

- Changing θ : This is how experiments explore physics
- Sensitivities & model uncertainties

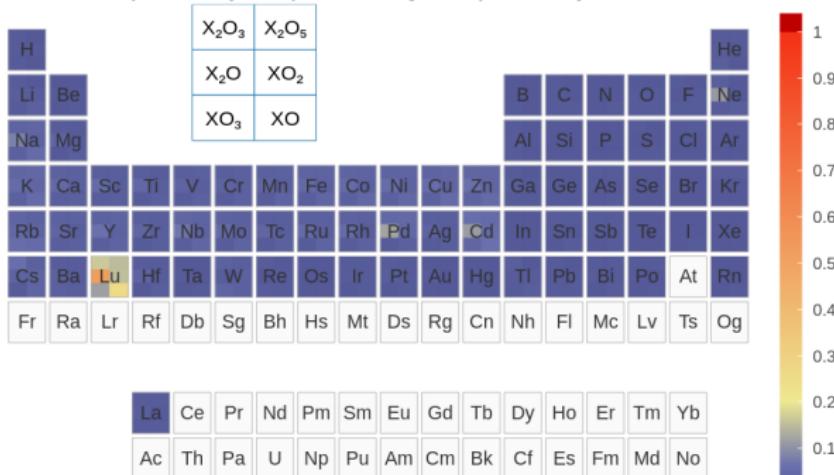
- A posteriori error: $x - x_* \approx -\nabla^2 E(x_*)^{-1} \nabla E(x)$

- Estimate accuracy of simulations

WIP: Integration into high-throughput frameworks

- Algorithms are only useful if they work in practice !
- DFTK plugin for AiiDA workflow manager
- Goal:** Automated testing of algorithms and error estimates
- Verification study Quantum-Espresso vs. DFTK
 - ⇒ Results agree, algorithms can outperform QE

ϵ for DFTK@PW|PseudoDojo-v0.5|rcut=10 vs. QE@PW|PseudoDojo-v0.5



Bruno Ploumhans

Contents

1 Example use cases for  DFTK

2 API and ecosystem integration

DEMO

DEMO



DFTK interface and ecosystem integration

Advertisement break

Open PostDoc in the EPFL ~~Mat Mat~~ group



Topic: Efficient inverse materials design

- Bayesian optimisation
- AD & gradient approaches
- Interdisciplinary environment of MARVEL:
Reproducible workflows, sustainable software,
computational materials discovery, statistical learning
- See <https://matmat.org/jobs/>



Psi-k workshop (*M. F. Herbst, A. Levitt, J. Haegeman*):

"Julia for numerical problems in quantum and solid-state physics"

- **26–28 November 2024** at EPFL, CECAM-HQ, Lausanne
 - Targets: Linear algebra, physics and julia communities
- ⇒ <https://www.cecams.org/workshop-details/1355> (Deadline: 20th Sep)

Summary and outlook

- Current state of  **DFTK**:
 - Unique robust material-adapting DFT algorithms
 - ForwardDiff to setup & solve response problems
 - Reduced settings (error analysis) and high-throughput testing
- Future work:
 - Explore error control & sensitivity (inverse design, surrogates)
 - Employ as frontend for domain-specific libraries (SIRIUS)
 - Composability with  **JuliaMolSim** (structure opt., surrogates, ...)
 - Bring methods to  **AiiDA** (for adoption and testing !)
- Where you can help:
 - Improve GPU performance (Hackaton anyone ?)
 - Parallelisation & performance bottle necks in AD / response
 - Explore alternative AD backends (Enzyme)
 - Use  **DFTK** &  **JuliaMolSim**, report bugs, enhance docs

Acknowledgements

MatMat group

- Bruno Ploumhans (~~MatMat~~)
- Anna Paulish (~~MatMat~~)
- Niklas Schmitz (~~MatMat~~)
- Cédric Travelletti (~~MatMat~~)

Aiida interface & verification

- Giovanni Pizzi (PSI)
- Junfeng Qiao (~~EPFL~~)
- Yihan Wu (~~EPFL~~)
- Austin Zadoks (~~EPFL~~)

Robust algorithms

- Eric Cancès (École des Ponts)
- Gaspard Kemlin (Université de Picardie)
- Antoine Levitt (Université Paris-Saclay)
- Benjamin Stamm (Stuttgart)
- Bonan Sun (~~EPFL~~)

- All > 40  **DFTK** contributors



Questions?

~~M~~~~X~~t Mat <https://matmat.org>

 mfherbst

 michael.herbst@epfl.ch



https://michael-herbst.com/talks/2024.10.21_JuliaMolSim_DFTK.pdf

 DFTK <https://dftk.org>